

# Time Warp Operating System

Parth Nain & Jay Lin

# Table of Contents

1. Virtual Time
  - a. Implementation
  - b. Time Scale Comparison, Leslie Lamport
2. Time Warp
  - a. Global Control Mechanism
  - b. Rollback definition
3. Time Warp Operating System (TWOS)
  - a. Discrete Event Simulations
  - b. Architecture
  - c. Comparisons with Typical O.S.
  - d. Rollback in TWOS
4. Summary

Papers discussed:

- "Distributed Simulation and the Time Warp Operating System" (David R. Jefferson et al)
- "Virtual Time" (David R. Jefferson)

# Virtual Time

- **Virtual time**: Proposed as a new paradigm for distributed computation, 1985
- **1D coordinate system** that extends globally (across all processors)
- **Values**:
  - $-\infty$ , the virtual time value in every system before computation begins
  - $\infty$ , all operations have been completed, and termination is ready
  - $(-\infty, \infty)$  = valid timestamps for events & messages
- Real time comes with various side effects for parallelized or distributed systems
- Advantage of virtual time is **total ordering and handling of simultaneous events**

# Virtual Time Implementation

- Numerous **local virtual clocks**, each hold own set of coordinates in the global system
- All primitive actions have both a coordinate in the time system, and a coordinate for **which processor** performed the action
  - Changing variables
  - Sending messages
- **Sender stamps messages are stamped with:**
  - Sender name
  - Virtual send time
  - Receiver name
  - Virtual receive time

# Virtual Time Implementation cont.

- Lamport's Clock Conditions

- Virtual send time of each message  $<$  its virtual receive time
- Virtual time of each event in a process  $<$  virtual time of the next event in the process
  - These conditions must be satisfied by all Virtual Time systems

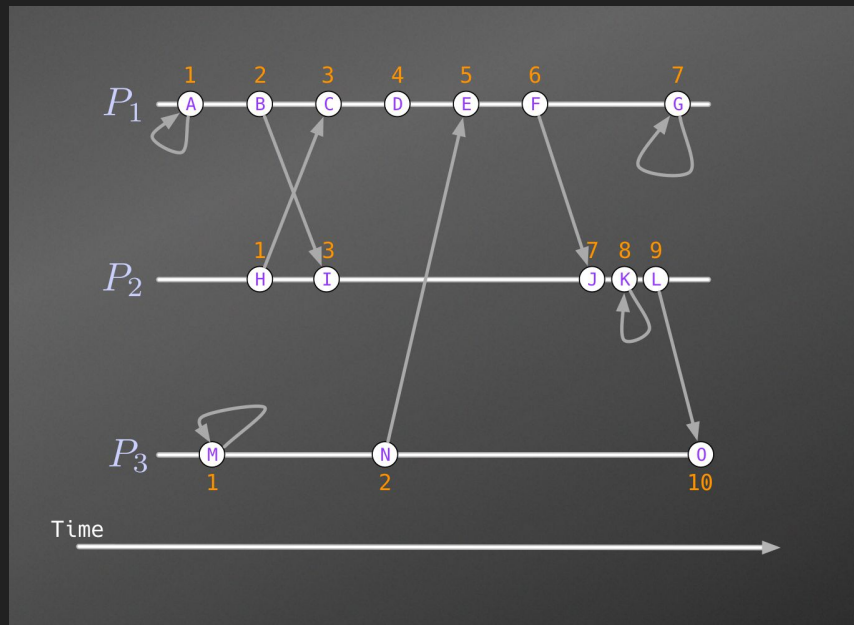
- Expected Event Operations

- Process may...
  - receive any # of messages stamped with their own location and timestamp, and read contents
  - read the virtual clock of the event
  - update its state variables
  - send any # of messages to other processes

# Time Scale Comparison

## Leslie Lamport

- Showed that simple relationships (before/after) only form partial order, and distinct events can become incomparable if concurrent in declaration
- Algorithm involves **labeling each event from a totally ordered set**
- Virtual time reverses this algorithm: **assumes that all events will be labeled**, and then leaves room to resolve problems by **rolling back later**



# Time Warp Mechanisms

## Local Control Mechanism



- Abstractly, **global time** is expected to **slowly progress** forward as a simulation **completes** (a singular concept)
- In implementation, **numerous local clocks with differing values** are interacting, still expected to move forward together
- Each process can receive messages and **order them by receive time**, and cycle through receiving-executing-sending. This will proceed unless a **misordered** message appears.

# Global Control Mechanism


Local clocks alone leave many problems unresolved...



- Is progress guaranteed with rollbacks happening?
  - Yes. Each local process regularly estimates GVT (Global Virtual Time), the lowest bound for all virtual times present in any sent message.
  - This is a commitment horizon that cannot be rolled back.
- How can all processes detect termination?
  - Termination is  $+\infty$ , when all timestamps reach  $+\infty$ , GVT will hit that value as well and termination will be signalled to all processes.
- Handling errors and output
  - Run-time errors mark a given state with an error flag, attempt to roll-back. Any errors timestamped before GVT will be presented after the simulation.
  - Output (ex: printing to a screen) is only performed when GVT passes the receive time of a message; cannot be rolled back.




# Rollback



**Undo** any consequence of an **incorrect action**, such as infinite loops and errors.

Rollback of a **message** = “**un-sending**”.


- **Anti-message**: message w/ negative (-) sign attached to it.
- **Annihilation**: anti-message is sent, and cancels out with the positive pair during a sender rollback.



# Rollback cont.

Message-antimessage pair per message process  $P \rightarrow Q$ .

- (+) goes to **receiving** process Q's **input** queue
- (-) goes to **sending** process P's **output** queue
  - Eventually **garbage-collected** if P doesn't rollback and timestamp < GVT (commitment time)
- If P rollback,
  - P's Local Virtual Time will point towards an earlier section of the **input/output queues**.
  - Any formerly emanated messages cancel/**annihilate** with the anti-message. May cause a chain of rollbacks



“Ripple effect” possibly caused by rollbacks

# Time Warp OS


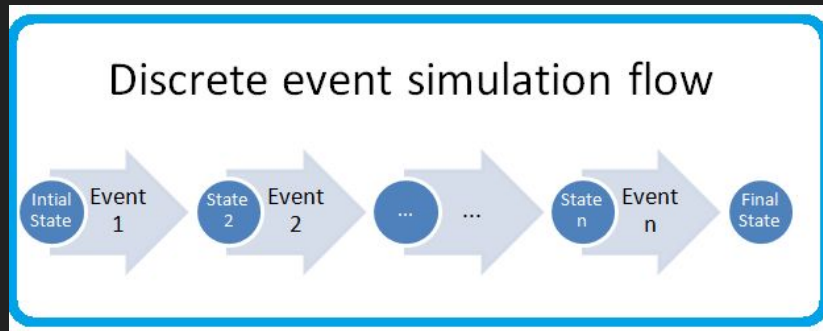
Developed in the 1980s in hopes of being able to concurrently run discrete-event simulations at an optimal speed.

- Different from typical OS; performs process synchronization using a general distributed process rollback mechanism.
- Written with C and assembly.

# Discrete Event Simulations (DES)

Simulates the **behavior of a real-life system** by modeling the system as a **sequence of events**, observing its **state changes in time**.

- Example real-life system: healthcare services, retail, military projects...
- Computationally \$\$\$, can take days to complete w/out parallelism.



# DES: How to speed them up?

1. Run different parts of the simulation in parallel, using a **time-stepped** approach.
  - a. Easy for **regular** systems that contain objects that change state at regular time intervals, like logical circuits with a clock.
  - b. **Very difficult to emulate systems that are not driven by regular time intervals**, or irregular.
2. **Event-triggered** approach
  - a. Should be used over the time-stepped method in the case of irregular systems as mentioned above, because time intervals are **irregular**.
  - b. TWOS comes into play as a solution



# Time Warp OS: background

Initially designed to be run on a hypercube, an arrangement of a **network of processors** running in **parallel**.

- TWOS can utilize any # of processors.

Why a **standalone** OS and not just on top of another OS?

- Rollback requires a **whole re-designing of operating system** scheduling, synchronization, flow control, memory management... etc.



Example diagram of a **hypercube** topology, each square is a node, or processor connected to a fixed # of neighbours

# Time Warp OS: Architecture

Top Layer = “**executive** layer”

- **C** code implementing Time Warp
- **Portable** (code not dependent on and can work on different platforms)

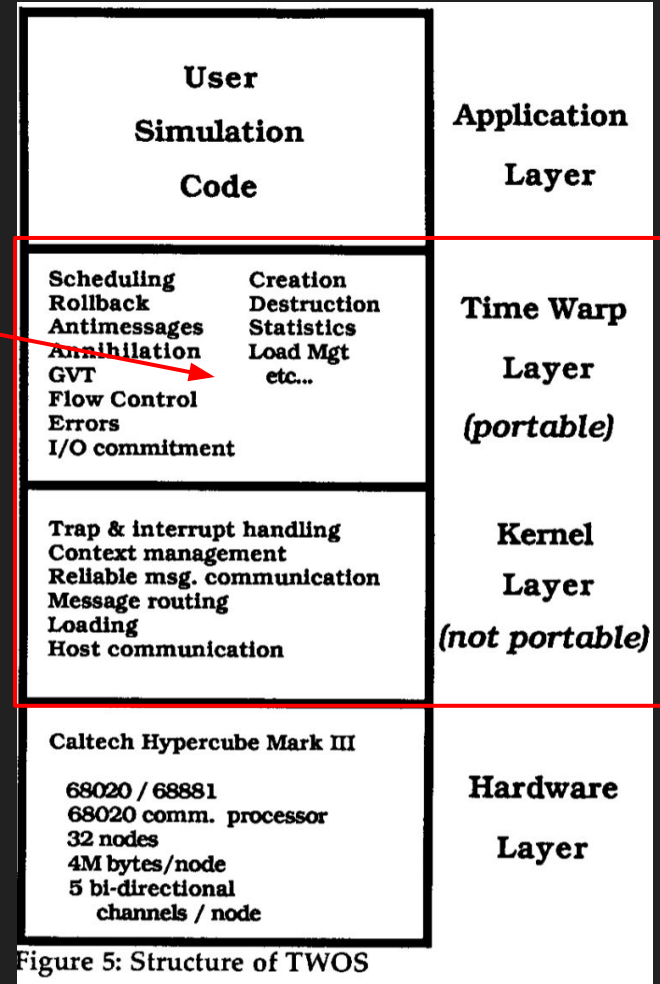


Figure 5: Structure of TWOS

# Differences b/t Time Warp and typical OS

Typical OS	TWOS
<ol style="list-style-type: none"><li>1. <b>FIFO</b> message queueing, preserves message sending order</li><li>2. Blocks sender in the case of <b>preventing overflow</b> of receiver's memory; easier storage management</li><li>3. <b>Channels</b> for processes to communicate</li><li>4. <b>Deletes message</b> as soon as receiver reads it to save memory</li></ol>	<ol style="list-style-type: none"><li>1. <b>Virtual timestamp order</b> message queuing</li><li>2. Difficult storage management; usually runs with <b>memory almost full</b></li><li>3. No channels; <b>any process send to any process</b></li><li>4. <b>Can't delete message</b> before <b>commitment time</b> (GVT); rollback might require it to be read again</li></ol>



# Processes in TWOS

Each **process** has an **input** queue, **output** queue, and event history (**snapshot** queue).


If a **process**  $Q$  receives a message from  $P$  whose timestamp  $t < \text{its LVT}$  (occurs in its past):

1.  $Q$  rollback its own event execution to that timestamp (look at snapshot queue)
2.  $Q$  also rollback its own input queue to that point
3. Send anti-messages to cancel (annihilate) all messages whose timestamp  $> t$

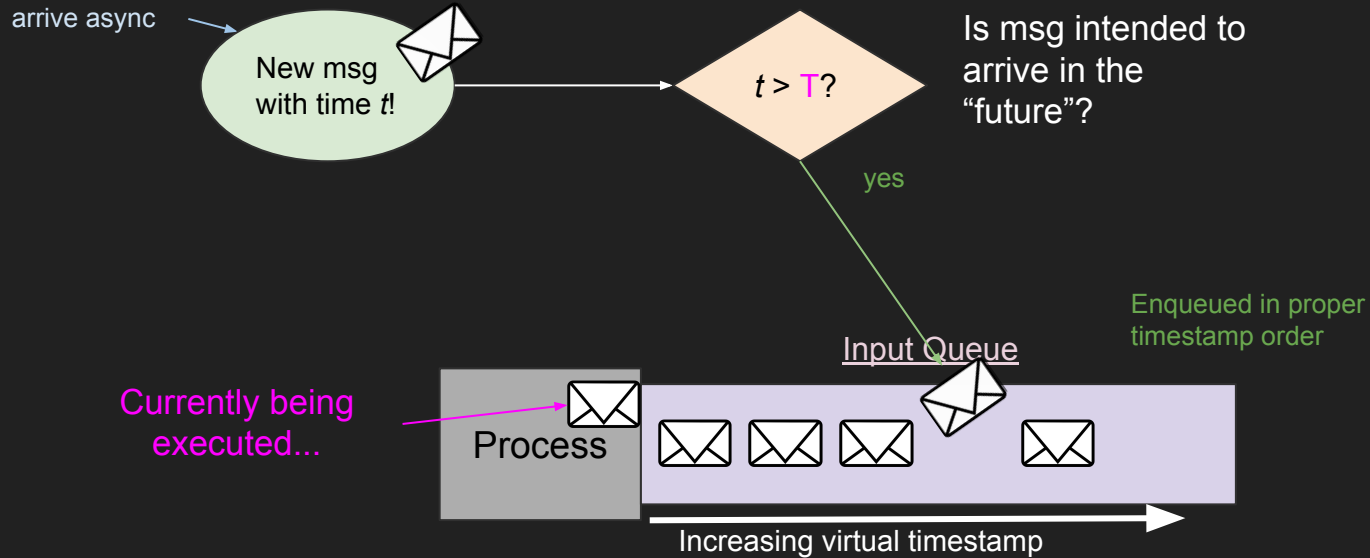


# Rollback in TWOS

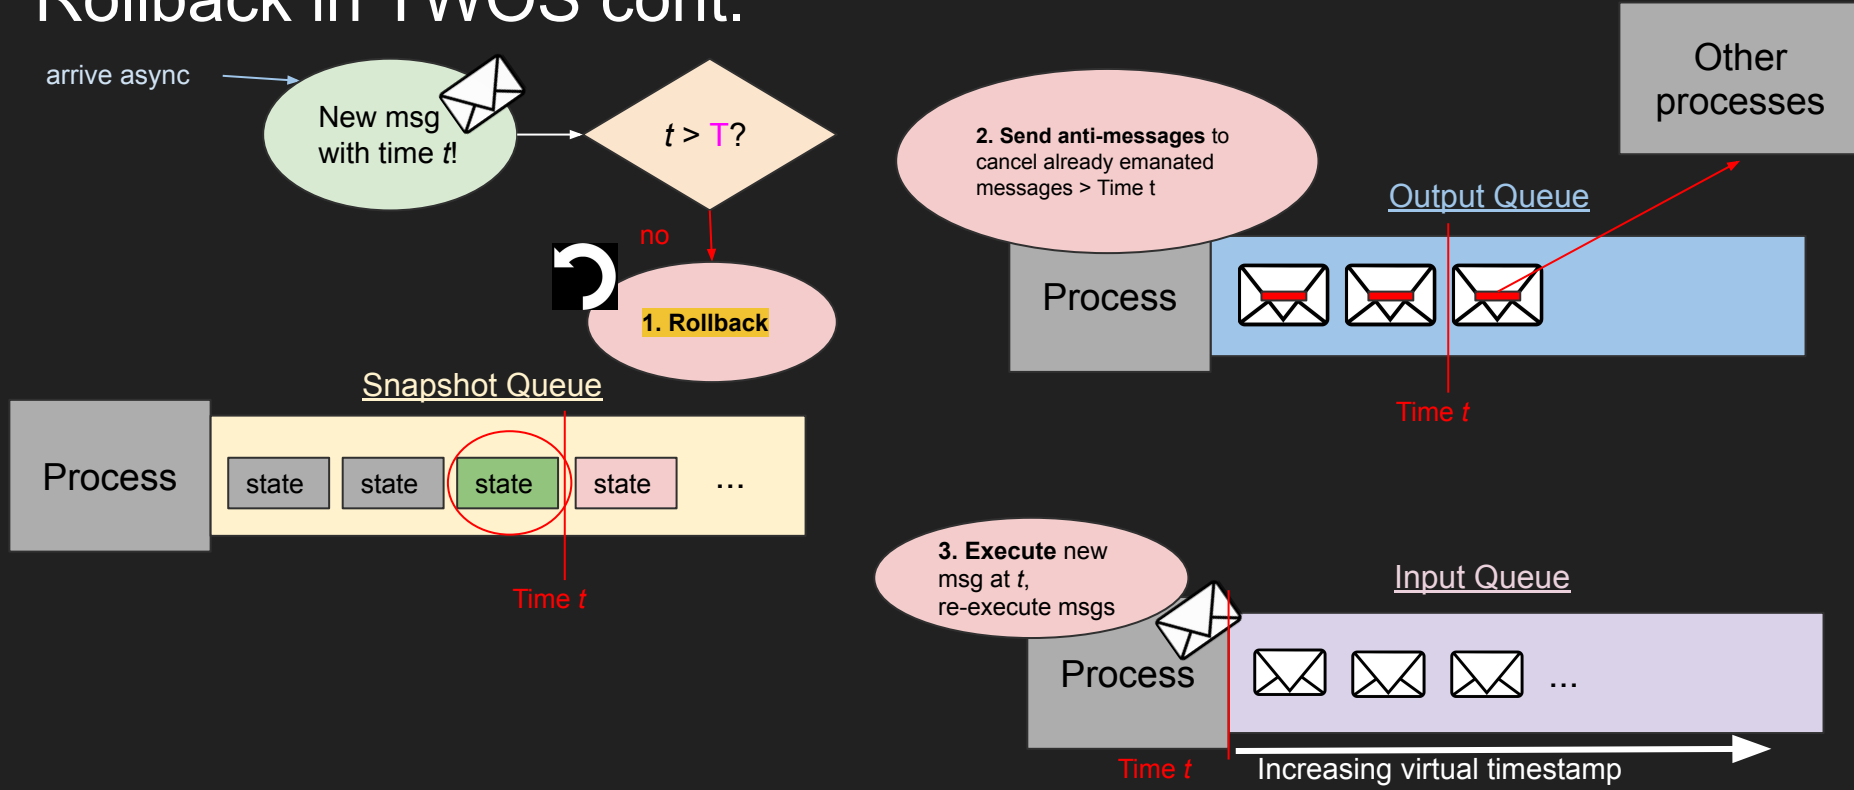
Optimistic synchronization: allows for temporary causal violations, assume that the messages in the input queue are the always true “next” messages to execute.




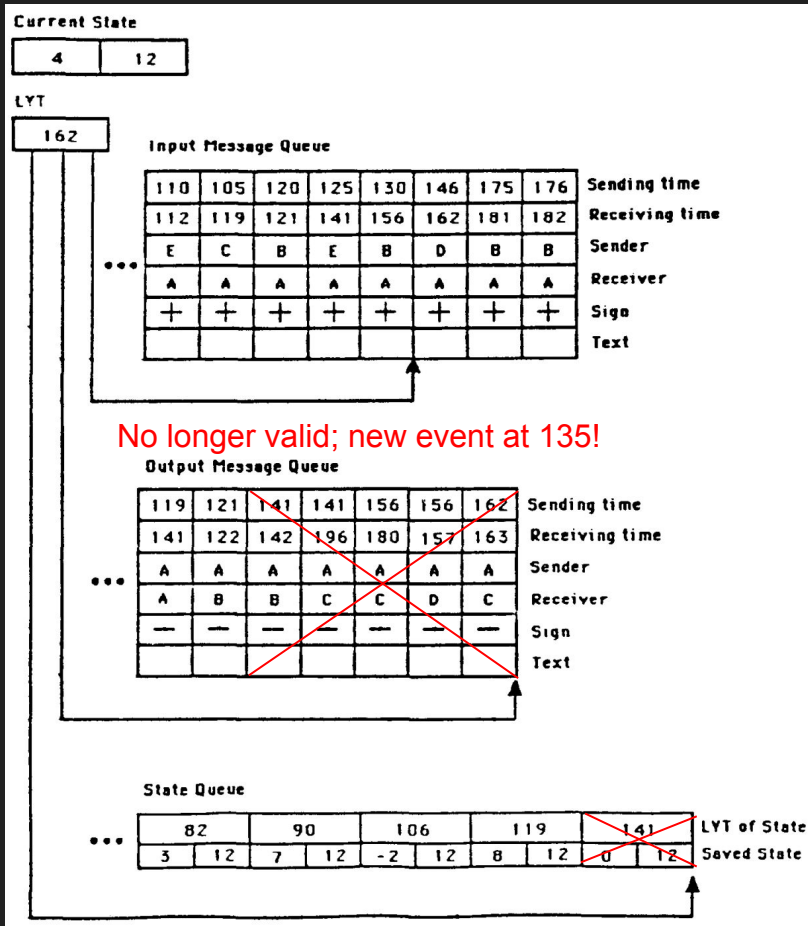
# Rollback in TWOS



# Rollback in TWOS cont.




# More detailed look at process queues and rollback




★ = LVT

# Timing Chart



★ = LVT

# Timing Chart



# TWOS impact and summary

Significant contribution to the speedup of running simulations in distributed systems.

- Speedup of simulations actually required more rollbacks contrary to initial beliefs.

Drawbacks of TWOS: high memory usage, huge overhead for rollbacks.

- Studies only done w/ limited # of processors (what happens on a larger scale?).
- Performance analysis experiments were only done on the Mark III computer (hypercube) in 1987; doesn't show us Time Warp's performance in other environments.

Personal remarks:

- TWOS should be tested on environments where memory is scarce/ on larger models (w/ more processors).
- Overall, the *Virtual Time* paper was much easier to digest because of its use of diagrams.



# References

1. David Jefferson, Brian Beckman, Fred Wieland, Leo Blume, Mike DiLoret, Phil Hontalas, Pierre Laroche, Kathy Sturdevant, Jack Tupman, Van Warren, John Wedel, Herb Younger, and Steve Bellonot, "**Distributed Simulation and the Time Warp Operating System**" Technical Report, UCLA, August, 1987.
2. Jefferson, D.R., "**Virtual Time**", ACM Transactions on Programming Languages and Systems, Vol.7 No.3, 1985, pages 404-425.
3. Stewart Robinson. 2004. *Simulation: The Practice of Model Development and Use*. John Wiley & Sons, Inc., Hoboken, NJ, USA.
4. <https://www.britannica.com/technology/hypercube>
5. [https://www.youtube.com/watch?v=OpHNPK6oSDs&ab\\_channel=DhananjaiRao](https://www.youtube.com/watch?v=OpHNPK6oSDs&ab_channel=DhananjaiRao)

Q&A